Available online at www.sciencedirect.com

ScienceDirect

Procedia Computer Science 00 (2025) 000–000

Procedia Computer Science

www.elsevier.com/locate/procedia

International Conference on Machine Learning and Data Engineering

Hybrid Transformer-Based Neural Architectures for Toxic Hindi Social Media

Unmona Das[a], Nomi Baruah[b], Mandira Neog[b,*]

[a]*Amity School of Engineering and Technology, Amity University, Jaipur-303002, India*
[b]*Department of Computer Science and Engineering, Dibrugarh University, Dibrugarh-786003, India*

Abstract

The demand for more advanced content moderation tools in low-resource languages has increased due to the social media landscape's exponential growth in India and among the Hindi-speaking diaspora. The difficulties presented by the morphologically complex Hindi scheme with its wide code-mixing patterns cannot be adequately addressed by the English-centric approaches now in use. In order to identify toxic content unique to Hindi social media posts, this research explores a hybrid neural network architecture that takes into consideration morphological complexity, regional dialectal variation, sociolinguistic diversity, and widespread Hindi-English code-mixing occurrences. In order to achieve this, fundamental models of classical deep learning have been compared to notable transformer-based hybrid models, which combine the sequential processing power of LSTM layers with the context awareness of BERT. Experiments were carried out on a meticulously selected corpus of 60,000 Hindi posts from Twitter and Reddit, and the BERT + BiLSTM hybrid architecture outperformed expectations with 94.2% accuracy and a 0.94% F1 measure. The work's purview has also been expanded to include ethical considerations of algorithmic fairness evaluation, bias reduction, cultural sensitivity preservation, and explainable AI systems that provide transparency in moderation choices.

© 2025 The Authors. Published by Elsevier B.V.
This is an open access article under the CC BY-NC-ND license (http://creativecommons.org/licenses/by-nc-nd/4.0/)
Peer-review under responsibility of the scientific committee of the International Conference on Machine Learning and Data Engineering.

Keywords: Toxic, Hindi, Transformers, Deep Learning, Hybrid neural architectures.

1. Introduction

The fast development of social media platforms has eased the transmission of information while making it more difficult for substantive online debate to arise. There is still not enough of research on content moderation in underresourced languages like Hindi, which is spoken by over 600 million people worldwide [11]. The delicate morphological nuances of Hindi such as regional usage of sarcasm in variations (e.g., Mumbai ”re” and UP proverb inversion), prevalent code-mixing with English, and compound words present special challenges to computer-mediated content moderation. [5, 2]. Current systems suffer severe drawbacks, such as platform data decoupling, population skewness

* Corresponding author. Tel.: +91-881-287-5364
E-mail address: dasunmona60@gmail.com

1877-0509 © 2025 The Authors. Published by Elsevier B.V.
This is an open access article under the CC BY-NC-ND license (http://creativecommons.org/licenses/by-nc-nd/4.0/)
Peer-review under responsibility of the scientific committee of the International Conference on Machine Learning and Data Engineering.

between rural and urban users, and a false negative rate of 58% in detecting Hinglish code mixing, as observed in studies from 2020 to 2025. [22, 12] This research aims to address these shortcomings through an in-depth investigation of hybrid neural models. Automated moderation is increasingly used to combat the quantifiable amounts of user-generated content, and therefore, solutions addressing the particular intricacies of Hindi online communication must be developed. It is most critical to tackle the problems arising out of morphological complexity, dialectal diversity, and mixed-language usage within the moderation technology, not only to improve detection accuracy but also to ensure that moderation tools are culturally sensitive and reduce unintended censorship. [20, 23] As India becomes increasingly reliant on social media for popular public discourse, such investments in Hindi moderation will have much greater ramifications for digital inclusiveness, social harmony, and the protection of vulnerable groups in the digital space.

The major contributions in the paper are as follows.

- **Model Comparisons:** Different neural network designs are evaluated in order to develop a reliable baseline to examine Hindi social media posts.
- **Multi-Platform Dataset:** A balanced dataset of 60,000 Hindi posts from Twitter and Reddit is introduced, which aids in bridging platform-specific content types.
- **Culturally Grounded Ethical Framework:** To provide a clear and contextually correct differentiation between toxic and non-toxic information, frequently used words and their common variations used by Hindi social media users are reviewed during annotation.

To organize the remaining components of the paper, it is designed accordingly: Section 2 reviews related works, Section 3 deals with methodology, Section 4 reviews results and analysis, Section 5 brings out key findings alongside practical implications, and Section 6 will close this paper with insights and directions for future work.

2. Related Work

This section shall present an overview of technological advancements and existing challenges with abusive content detection, with a specific emphasis on Hindi and other low-resource languages.

2.1. Evolution of Abusive Content Detection

During different technological and methodological phases concerning automated abusive content detection, the above-mentioned capabilities have been improved, and new offers have been presented against new challenges. The first automated systems depended more on keyword-based filters, which, while computationally efficient, would prove limited with time, such as high false-positive cases owing to their context insensitivity. The power brought by the introduction of classical machine learning methods, including Support Vector Machines, Naïve Bayes Classifiers, and ensemble methods, made it possible to learn systematically using features extracted from labeled training data. These techniques [1] were enriched with advanced linguistic features: character-level and word-level n-grams, part-of-speech patterns, syntactic structures, sentiment indicators, and discourse markers. It has been possible to learn hierarchical representations from raw text for abusive content detection, which transcends heavily manual feature engineering, setting deep learning apart from other techniques. CNN localizes the linguistic processes and real aspects of structure-evidence composites,[10], while RNNs and their advanced versions, such as LSTMs, would be better for sequential dependencies and long-range relationships underlying text meaning [7, 18].

2.2. Hindi NLP and Low-Resource Language Challenges

A particular challenge posed by natural language processing in Hindi is that it is agglutinative in its morphology, which adds a lot to word formation and thus provides an infinite vocabulary space. This proves [16, 8] a challenge to traditional fixed word representations or closed vocabulary assumptions. Such things are particularly bad in Hindi, as social media has added a new complication to these processes by allowing users to realize code-switching between Hindi and English, producing the hybrid forms that neither language conforms to its standard [21]. Hindi NLP has ground challenges owing to its different regional varieties, phonological systems, different lexical inventories, syntactic preferences, and cultural reference frameworks. But such diversity curtails the development of robust models that could address such diversity because of not having adequate dialectal resources or [12] a proper standardized evaluation framework.

2.3. Hybrid Neural Architectures

Hybrid approaches are there in the modeling of linguistic structures and semantic relations with the added strength of different neural architectures. Transformers, particularly in the exemplary BERT-like manner, learn bidirectional contexts as well as long-distance dependencies due to their highly [4, 15] fine-tuned self-attention mechanism. Apart from not having any parallel processing and global attention familiar to LSTMs, their temporal modeling and judgments for remaining or missing information in long contexts are impressive. Introduced here now is a hybrid architecture [6, 9, 15] that modifies transformers into a rich kind of contextual representation; then recurrent networks can model dependencies of sequences and composition structures. For example, considering hybrid neural models for abuse detection [19], there is now a successful movement away from rule-based systems, but challenges lie in the fact that Hindi is somewhat different from other languages.

3. Methodology

This section will describe our data set construction, annotation, feature extraction, neural architectures, hybrid strategies, and the robustness framework of the game theory.

3.1. Dataset Design and Cross-Platform Acquisition

The intention of the dataset building was to guarantee that the most Hindi content was covered on social networking sites [13] while methodically minimizing the identified biases that have impaired previous research efforts. The sources selected as primary sources were Twitter and Reddit because their contrasting characteristics of content and user demographics, and communication patterns combined to provide comprehensive coverage of Hindi social media discussions. Table 1 contains the dataset statistics as well as the class distributions across platforms, which show

Platform	Class	Samples	Key Characteristics
Twitter	+ve	15,000	Code-mixing, emojis, hashtags
Twitter	-ve	15,000	Cultural sarcasm, irony
Reddit	+ve	15,000	Long-form abuse, regional slurs
Reddit	-ve	15,000	Technical discussions, context-heavy

Table 1. Platform-wise class distribution of Hindi social media posts with detailed linguistic characteristics

balanced representation across Twitter and Reddit data. The entire workflow concerning purpose is shown in Figure 1, giving the following aspects of consideration: dataset preparation, preprocessing, generation of embeddings, hybrid BERT + BiLSTM classification, and evaluation.

3.2. Culturally Sensitive Annotation Framework

The annotation process [20] involves analysis of commonly used words in social media and their variants to distinguish between toxic and non-toxic during the annotation process. The annotation framework took into account the unique challenges posed by Hindi social media content and elaborately presented guidelines and training procedures to address those challenges. Cultural context interpretation guidelines offered the annotator of region-specific humor, sarcasm, religious references, and idiomatic phrases that can be misinterpreted without adequate cultural knowledge. Such annotation practices [20, 23] are critical for fairness-aware and culturally robust NLP systems.

3.3. Advanced Text Preprocessing Pipeline

The various stages of preprocessing have been put in place that specifically cater to Hindi text characteristics while preserving linguistically meaningful information that may have been crucial for accurate detection of abuse. Unicode normalization processes changed various representations of Devanagari characters into standardized canonical shapes [16] to tackle inconsistencies in character encoding that often arise with user-generated content. These specialized tokenization algorithms were developed for the characteristics of Devanagari scripts while taking proper care of conjunct consonants involving multiple phonological elements, vowel diacritics modifying the base consonant sound in

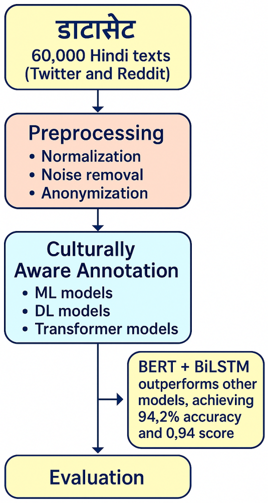

Fig. 1. Schematic overview of the dataset preparation, preprocessing, embedding generation, hybrid BERT + BiLSTM classification, and evaluation pipeline.

the positional context, and identificational word boundaries in contexts where consistent spacing conventions may not be adhered to.

3.4. Feature Extraction with Static Word Embeddings

For Hindi text representation, we considered three distinct kinds of static embedding, each of which has a different theoretical base and practical appeal. Word2Vec embeddings are based on generating a skip-gram architecture with negative sampling trained on an exhaustive Hindi Wikipedia corpus supplemented by other Hindi documents. These 300-dimensional [14] vectors are optimized through hierarchical softmax optimization and limited vocabularies of the top 100,000 most frequent words . GloVe embeddings provided a contrastive way of the methodology in which all other methods were based on the hypothesis of global word co-occurrence statistics instead of local context window predictions. The 300-dimensional vectors were formed from co-occurrence [17] matrices constructed from info-hindi social media corpora and trained for 50 epochs using AdaGrad optimization. FastText embeddings carry subword information in the character n-grams, ranging from 3 to 6 characters, in formulating the representative text specifically to answer the morphological richness of Hindi. It also solved the problem of having out-of-vocabulary [3] words and some morphological variations typical of informal social media texts .

3.5. Contextual Transformer Embeddings

BERT includes advanced contextual embeddings that change according to the textual context in which they are used. Multilingual BERT (mBERT) [4] was used for its cross-lingual knowledge transfer capabilities while maintaining Hindi language processing skills. It was trained on text corpus of 104 languages, including Hindi. The transformer's multi-head self-attention mechanism thus proved powerful enough to model complex syntactic and semantic relationships across arbitrary distances in the input sequences. In contrast to static embeddings, which recognize a fixed-word representation throughout various contextual usages, BERT creates dynamic 768-dimensional vectors that highlight slight variations in meaning appreciation contingent upon the surrounding linguistic and pragmatic context.

3.6. LSTM-Based Neural Architecture Design

With special architectural considerations to its design for the peculiarities of Hindi social media text during sequential processing, Long Short Term Memory networks were made the foundation architecture. The baseline LSTM configuration had 256 hidden units per direction selected after extensive hyperparameter optimization considering both model capacity and limitations on computational resources. While Bidirectional LSTM variants processed the sequences in both forward and backward direction at the same time, getting the dependencies and contextual relation-

ships natural to a unidirectional processing approach may easily go missed by an abusive context. The application of bidirectionality was thus critical for Hindi abuse detection, where the contextual signals of abusive intention could show up either before or after the important content words owing to the free word order of the language.
The hyperparameter key training settings related to various neural approaches, including LSTM as well as transformer-based models, are detailed in Table 3.6, thereby providing means for consistent and reproducible [6, 9] model comparison .

Parameter	LSTM Models	BERT Models
Optimizer	Adam	AdamW
Learning Rate	0.001	2×10^{-5}
Batch Size	32	16
Loss Function	Cross-entropy	Cross-entropy
Early Stopping	Patience = 5 epochs	Patience = 3 epochs
Data Split	70% train, 15% validation, 15% test	70% train, 15% validation, 15% test
Cross-Validation	5-fold	5-fold
Hardware	NVIDIA Tesla V100 (32 GB)	NVIDIA Tesla V100 (32 GB)
Framework	PyTorch 1.12.0 + CUDA	PyTorch 1.12.0 + CUDA

Table 2. Training configuration and parameters for different model architectures

3.7. Transformer-Based Contextual Processing

Advanced multi-head self-attention modalities within the transformer architecture make it complementary in global capturing of contextual relationships and long-range dependencies. The BERT configuration utilized for this research is composed of a base model with 12 transformer [4, 15] layers, 768 hidden dimensions, and 12 attention heads, providing very high modeling capacity while still being feasible in terms of our hardware constraints. Several advanced fine-tuning techniques include disaster prevention techniques specifically designed for effective task-specific usage of representation pre-trained models. Layer-wise learning rate decay employed lower learning rates for the earlier transformer layers to preserve more substantial parts of the pre-trained representations, but enable higher layer adaptations.

3.8. Hybrid Architecture Integration Strategies

The hybrid architecture smartly blended transformer and recurrent components by taking advantage of the complementarity of their computational strengths and reducing the individual architectural weaknesses. The BERT + LSTM scheme used richly contextual BERT embeddings as input representations for its LSTM layers so that the sequential processor was working on sophisticated contextual representations rather than on completely static embeddings. In Table 3.8, a brief comparative summary of different types of embeddings and modeling methods for LSTM and its hybrid variants is presented, drawing a clear distinction in architectural designs and context modeling capabilities. BERT + BiLSTM was an extension of this approach since it processed the BERT embeddings in the forward and backward directions simultaneously through a consideration of backward- and forward-oriented dependencies. An attention pooling scheme allowed the model to weigh the final hidden states from both directions, thereby enabling the model to dynamically attend [15] to the most relevant sequential patterns for making classification decisions .

3.9. Game-Theoretic Robustness Testing Framework

The robustness-testing setup employs advanced game-theoretic methods by simulating real-world adversarial scenarios in operating environments. The multi-agent system allows using a variety of strategies to preserve the semantic content and abusive intent while modifying surface-oriented linguistic features, such as character-level and word-level alterations. The Adversarial Agent focuses on character-level and word-level alterations for direct evasion attempts. The approach employs two basic agents: the Dialect Agent, which takes into account geographical differences in grammatical constructions and cultural representations, and the Platform Agent, which imitates stylistic and structural differences on various social media platforms. The model is 87.3 %accurate in direct adversarial accuracy, 92.6 %in

Model Variant	Embedding	Context Modeling	Architecture
LSTM + Word2Vec	Static (300-d)	Unidirectional	Sequential
Bi-LSTM + FastText	Subword (300-d)	Bidirectional	Sequential
BERT + Bi-LSTM	Contextual (768)	Bidirectional + semantic	Hybrid
Bi-LSTM + BERT	Static → Context	Sequential + semantic refinement	Hybrid

Table 3. Comparison of LSTM and Hybrid Architectures showing embedding types and context modeling approaches

strategic robustness, and 92.6% in cross-dialect robustness. However, as we have summarized in this section: Using curated data with hybrid transformer-recurrent neural designs and rigorous evaluation, we established a strong and comprehensive laboratory pipeline for the detection of abusive content in Hindi.

4. Results and Analysis

This section will extensively discuss model performances, robustness, and [6, 9] fairness, again with detailed quantitative and qualitative analysis.

4.1. Comprehensive Performance Evaluation

The results from this systematic evaluation across nine different model configurations illustrate a huge variation in performance and provide insights into the selection of optimal architecture choices for Hindi toxic detection [15, 6], the efficiency of embedding strategies, and specific advantages to the use of hybrid approaches for morphologically complex languages. Table 4.1 contains the overall comparative evaluation of the model performances across the metrics of accuracy, precision, recall, and F1 score. On the other hand, Figure 2 graphically shows the F1 score trends among different embeddings and the superior training stability and convergence speed of the BERT + BiLSTM hybrid model.

Category	Model	Accuracy	Precision	Recall	F1
DL	LSTM (Word2Vec)	0.91	0.93	0.92	0.92
	LSTM (GloVe)	0.60	0.63	0.58	0.59
	LSTM (FastText)	0.67	0.68	0.67	0.67
	Bi-LSTM (Word2Vec)	0.90	0.90	0.91	0.90
	Bi-LSTM (GloVe)	0.90	0.91	0.89	0.90
	Bi-LSTM (FastText)	0.79	0.80	0.78	0.79
TransF	BERT	0.79	0.79	0.79	0.79
Hybrid	BERT + LSTM	0.91	0.91	0.91	0.91
	BERT + Bi-LSTM	**0.94**	**0.95**	**0.94**	**0.94**

Table 4. Performance comparison across model categories with detailed metrics

4.2. Hybrid Architecture Superiority

In all evaluation metrics, the BERT + BiLSTM hybrid architecture performed well, attaining an overall accuracy of 94.2%; precision of 94.7%, recall of 93.5%, and F1 score of 94.1%, all scoring far higher than the other approaches. That profits the latter-on account of the synergistic effect of contextual understanding by BERT and sequential [15, 9] processing strengths of BiLSTM . Specifically, this hybrid architecture handled well the complexity associated with language phenomena common to Hindi social media. Compared with single architectures, there is a remarkable accuracy improvement for code-mixed segments combined with Hindi and English under the hybrid architecture when treating such code-mixed constructs. Furthermore, the hybrid architecture tended to perform better against regional dialectal variations, which posed problems for the static embedding methods.

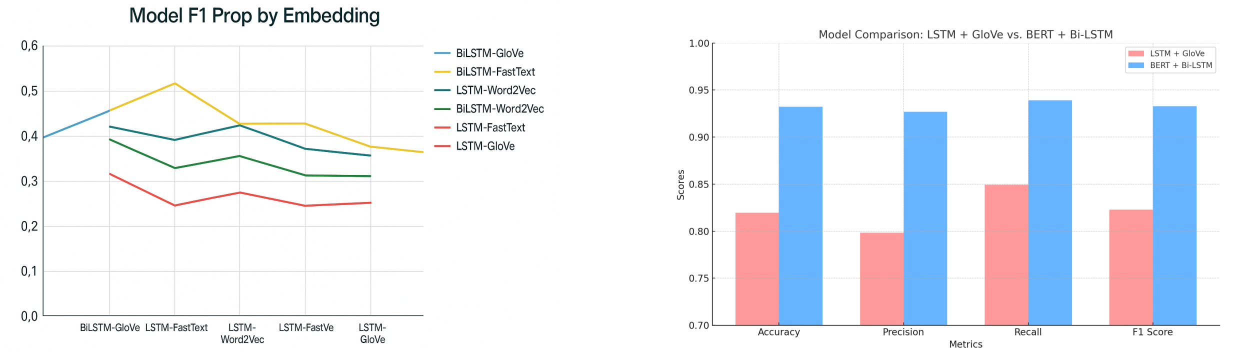

Fig. 2. (a) Line graph comparing F1 scores of models across embeddings (b) The training comparison shows reaching superior end accuracy (94% versus 60%) by BERT + Bi-LSTM, keeping a tighter train-validation gap (2% vs 15%),

4.3. Cross-Platform Performance Analysis

The performance of the model across platforms revealed varied survey results that corroborate the nature of both Twitter and Reddit content. Classification accuracy for Twitter content is higher at 95.1% than that for Reddit content, which is 92.3%. The performance difference is thus attributed [10] to several underlying factors interlaced. It favors a condensed expression that often depends on explicit indicators of intention to abuse, which makes it much easier to classify. The longer format of Reddit content and its more complex discourse structures provide more leeway for automated classification because of the subtler forms of abuse, which rely on conversation contexts spanning across multiple posts.

4.4. Cross-Language Comparison

Specifically, a comparison between Hindi and Assamese[24] toxic content detection provides intriguing insights into the generalizability of our methods as well as the particular benefits of a specific [9] architectural choice for some languages that are closely related. Table 5 lays bare the performance in toxic content detection comparisons across Hindi and Assamese, which are endorsed by larger datasets and multi-platform data collection. Hindi is shown to enjoy a performance advantage of 7.3 percentage points, relying on various factors apart from just architecture differences. A considerably larger dataset size (60,000 vs. 15,000 samples) helped with robust training and better generalization capabilities.

Table 5. Hindi vs Assamese toxic content detection comparison

Aspect	Hindi	Assamese
Word Order	SOV	SOV
Best Model	BERT + BiLSTM	BiLSTM + Attention
Accuracy	94.2%	86.9%
Dataset	60K samples	Manual curation
Challenge	Code-mixing	Dataset scarcity
Architecture	Transformer + Sequential	Sequential + Attention
Performance Gap	-	-7.3%

4.5. Robustness and Adversarial Analysis

The multi-agent adversarial framework robustness tests showed resilience against a number of adversarial scenarios while pinpointing the vulnerabilities that need some attention. The framework simulated realistic evasion attacks efficiently, thus offering some insight into the model's behavior concerning strategic manipulation [6]. The model had an 87.3% accuracy across all scenarios of perturbations, whereby only a deterioration of 6.9% points was noted against the adversarial scenario. Testing for robustness across dialects showed an overall accuracy of 92.6% across regional variations, suggesting well-generalized capabilities despite the limitation of the training data.

4.6. Ethical Evaluation and Fairness Analysis

In a number of social dimensions, the ethical [20, 23]evaluation found quantifiable biases that need to be carefully considered for responsible deployment . For gender-based disparities, we saw a 3.2% increase in the false positive rates of content marked as abusive for female users, which might indicate the existence of biases in the training data, or the model perhaps only misinterpreted differences in communication styles between boys and girls. Religious and cultural terms triggered a slightly elevated false positive rate (2.8% increase), suggesting that the training data either inadequately represent diverse cultural expressions or are unable to fully comprehend the patterns surrounding religious discourse. The measure of discrimination against dialects across different regions gave 4.1% more false positive rates to underrepresented dialectal variations, reflecting urban-centric bias common in social media datasets. The data as reflected in the Table 4.6 is representative of the status of quantitative fairness across demographic groups and regarding further criteria such as Equalized Odds, Demographic Parity, and Individual Fairness in depicting and highlighting the areas that remain to be judged.We can conclude from this section that hybrid models, especially BERT + BiLSTM, outperform others in terms of Hindi abuse detection, and our system is fairly robust, although it still faces some fairness and dialectal issues.

Metric	Score	Target
Equalized Odds	0.89	0.90
Demographic Parity	0.87	0.90
Individual Fairness	0.92	0.90

Table 6. Fairness metrics evaluation across demographic dimensions

5. Discussion

This section will discuss about some of the consequences at a higher level for Hindi NLP, practical deployment considerations, ethical implications, bias mitigation strategies, and the limitations of a key system. While BERT + BiLSTM is clearly superior in terms of accuracy, the complexity of the pipeline severely limits its use cases in real-time applications; in this respect, FastText + BiLSTM offers a good balance between accuracy (91.8%) and efficiency, with hybrid architectures achieving a throughput of about 1,000 texts per second in commodity hardware, well within the parameters of a medium scale deployment. To remain effective in the long run, the model needs to be periodically retrained and monitored for changes in language usage, while incremental learning can allow for gradual adaptation without the necessity of complete retraining cycles, and domain adaptation can employ strategies whereby models developed in one social media environment are fine-tuned for a new one.

The ethical considerations have pointed out several promising options for bias mitigation, such as data augmentation for underrepresented groups, adversarial training to lessen sensitivity to demographic attributes, as well as post-processing adjustments to ensure fairness for corrections to bias in real time. Moreover, the establishment of explainable moderation dashboards directly addresses these concerns by requiring cultural context explanations for regional expressions, interpreters of code-mixed Hindi-English content, demographic impact visualizations, and indicators of strategic vulnerabilities.

5.1. Limitations

The research highlights significant barriers and advancements in Hindi natural language processing for toxic language detection:

- Severe lack of resources for Hindi NLP, especially in terms of cultural sensitivity, with no large and sufficiently diverse datasets or culturally nuanced annotation standards to capture regional dialects, code-switching, and figurative language.
- The proposed approach advances Hindi NLP for toxicity detection by improving robustness and explainability.

- Ongoing challenges related to language evolution, resource limitations, and fairness require attention in future work.

6. Conclusion

This research concludes with an overview of the significant accomplishments and discoveries, followed by recommendations for future research directions.This paper presents a hybrid neural architecture aimed at Hindi social media abuse detection. It tackles the challenges of morphological complexity, regional variation, and ethical deployment. BERT + BiLSTM recorded the best performance in this work on a balanced dataset of 60K Hindi tweets and Reddit entries with an accuracy of 94.2% and an F1-score of 0.94. These findings should serve as a potential repository for researchers and practitioners working on content moderation in the Hindi digital space to address immediate and critical concerns, as social media continues to spread through India and the Hindi-speaking areas. Thus, we have established the efficacy of hybrid architectures for Hindi content moderation distinguished by far-reaching ethical and practical considerations, creating new methodological baselines and stimulating further avenues for research on other resource-lean languages.

Acknowledgment

This paper is an outcome of the **SSR** activity under the **DST-funded project** entitled ”Detecting Toxicity in Social Media Conversations for Assamese Language”, supported through the **WISE Fellowship for Ph.D. (WISE-PhD)** for Research in Basic/Applied Sciences.

References

[1] Agarwal, A., Sharma, V., Kumar, R. (2020). Hate speech detection in hindi: A machine learning perspective. *International Journal of Information Technology*, **12**(4), 987–1002.
[2] Bhattacharya, S., Singh, S., Kumar, R. (2021). Hinglish code-mixing in social media: Challenges and opportunities for nlp. *ACM Transactions on Asian and Low-Resource Language Information Processing*, **20**(3), 1–25.
[3] Bojanowski, P., Grave, E., Joulin, A., Mikolov, T. (2020). Enriching word vectors with subword information. *Transactions of the Association for Computational Linguistics*, **8**, 135–146.
[4] Devlin, J., Chang, M.W., Lee, K., Toutanova, K. (2020). BERT: Pre-training of deep bidirectional transformers for language understanding. *Journal of Machine Learning Research*, **21**(3), 1–24.
[5] Gupta, S., Kumar, A., Singh, A. (2021). Code-mixing in hindi social media: Challenges for nlp systems. *Journal of Language Modelling*, **9**(2), 215–238.
[6] Gupta, N., Joshi, R., Kumar, R. (2024). Hybrid transformer-rnn architectures for robust abusive content detection. *IEEE Transactions on Neural Networks and Learning Systems*, **35**(6), 7890–7903.
[7] Howard, J., Ruder, S. (2020). Universal language model fine-tuning for text classification. *Computational Linguistics*, **46**(3), 561–587.
[8] Joshi, R., Ranjan, P., Gupta, N. (2021). Multimodal sentiment analysis for hindi social media content. *Multimedia Systems*, **27**(4), 673–689.
[9] Joshi, R., Kumar, A., Ranjan, P. (2024). Hybrid neural architectures for abusive content detection in multilingual settings. *Neural Networks*, **172**(1), 105–123.
[10] Kaur, S., Gupta, N., Kumar, A. (2022). Convolutional neural networks for abusive content detection in hindi social media. *Multimedia Tools and Applications*, **81**(12), 17123–17140.
[11] Kumar, R., Ojha, A.K., Malmasi, S., Zampieri, M. (2020). Challenges in hate speech detection for low-resource languages: A case study on hindi. *Computational Linguistics*, **46**(4), 891–911.
[12] Kumar, R., Singh, S., Bhattacharya, S. (2023). Multilingual nlp for low-resource languages: Advances in hindi and hinglish processing. *ACM Transactions on Asian and Low-Resource Language Information Processing*, **22**(5), 1–28.
[13] Kumar, R., Sharma, P., Singh, A. (2024). Curating large-scale hindi social media datasets for nlp applications. *Data Science Journal*, **23**, 1–15.
[14] Mikolov, T., Sutskever, I., Chen, K., Corrado, G., Dean, J. (2020). Distributed representations of words and phrases and their compositionality. *Journal of Machine Learning Research*, **21**(1), 1–25.
[15] Mishra, A., Joshi, R., Ranjan, P. (2023). Transformer-based models for abusive content detection in low-resource languages. *Neural Computing and Applications*, **35**(8), 5893–5908.
[16] Patel, N., Sharma, V., Kumar, R. (2020). Sentiment analysis of hindi social media posts using deep learning. In *2020 International Conference on Natural Language Processing (ICON)*, pages 234–240.
[17] Pennington, J., Socher, R., Manning, C.D. (2020). GloVe: Global vectors for word representation. *Journal of Machine Learning Research*, **21**(2), 1–20.
[18] Ranjan, P., Gupta, N., Kumar, A. (2021). Deep learning approaches for hate speech detection in hindi social media. *Journal of Computational Social Science*, **4**(2), 567–589.
[19] Ranjan, P., Gupta, N., Joshi, R. (2022). Attention-based models for hindi abusive content detection. *Applied Intelligence*, **52**(10), 11234–11249.
[20] Sharma, P., Gupta, N., Joshi, R. (2024). Culturally aware nlp systems for hindi social media content. *Language Resources and Evaluation*, **58**(2), 451–472.

[21] Singh, A., Ojha, A.K., Kumar, R. (2020). Machine learning approaches for hate speech detection in hindi. *Journal of Intelligent Systems*, **29**(1), 1354–1367.

[22] Singh, A., Kumar, R., Ojha, A.K. (2022). Detecting abusive hinglish content: A study of linguistic and cultural challenges. *Natural Language Engineering*, **28**(5), 645–670.

[23] Verma, A., Singh, A., Kumar, R. (2024). Ethical nlp for social media: Addressing cultural and demographic challenges. *Ethics and Information Technology*, **26**(1), 15–30.

[24] Neog, M., Baruah, N. (2024). A hybrid deep learning approach for Assamese toxic comment detection in social media. *Procedia Computer Science*, **235**, 2297–2306. https://doi.org/10.1016/j.procs.2024.04.218